A Legofit Tutorial

Alan R. Rogers

July 9, 2022

Contents

1 Introduction

This is a tutorial on Legofit, a package that uses genetic data to estimate the history of population size, subdivision, and admixture [8, 9]. It will step you through the process of installing the software, using it to tabulate the frequencies of nucleotide site patterns, and then analyzing these data to fit a model of population history. We will assume that data files are already available in .vcf format. The full Legofit documentation is available online.

This tutorial will be used in the 2022 edition of the AGAR Workshop. During the workshop, we will do only last step in this process, which estimates population history from site pattern frequencies. We will not do the preliminary steps, which tabulate these frequencies. Nonetheless, this tutorial describes the whole process.

We plan to introduce students to the Center for High-Performance Computing (CHPC) at the University of Utah, and some of the sections below assume that you have an account on that system. (We'll be getting accounts for everyone enrolled in the workshop.) These sections involve using the SLURM Workload Manager to submit jobs to the compute cluster. However, our focus during the workshop will be on the last step, which can be done in two ways: either on the CHPC (using SLURM), or on your own computer (using the bash shell). This tutorial will cover both environments.

The material below is organized into the following sections:

2 Installing Legofit on your own machine

3 Setting up your environment at the CHPC

4 Making an input file in .raf format

5 Tabulating site pattern frequencies

6 Fitting models of history

7 Model selection and model averaging

8 Graphing results

Try to install Legofit (sec. 2) on your own before the workshop starts. We'll help with installation problems during the workshop, but that will take less time if you've already made a start. The workshop will not cover secs. 4–5. Those sections are only for reference, when you work with Legofit on your own. During the workshop, most of our efforts will be on fitting models (sec. 6), but we'll spend some time at the end on model selection, model averaging, and graphing (secs. 7–8).

2   Installing Legofit

This section explains how to get Legofit running on your own computer. It's already installed on the CHPC, so these steps are not needed there.

2.1   Install git if necessary

You may have git already. To find out whether you do, type:

```
type git
```

On Mac or Linux, this will print the location of the executable git file. If there is no such file, you'll get an error message saying `type:  git:  not found`. I recommend using homebrew to install packages on a Mac. To install git using homebrew, type:

```
brew install git
```

Otherwise, visit git's website and look for the download instructions.

2.2 Clone legofit

Use git to clone Legofit onto your machine. This step will create a subdirectory called "legofit," which will contain the source code. Before executing the command below, use the `cd` command to move into the directory where you keep source code that other people have written. I keep such code in a directory called `distrib`, which I originally created by typing `mkdir distrib`. Having created such a directory, move into it by typing

```
cd distrib
```

Then clone legofit by typing:

```
git clone https://github.com/alanrogers/legofit.git legofit
```

This will create a directory called legofit. Use `cd` to move into it.

2.3 Install a C compiler if necessary

To see whether you have a C compiler already, type `type cc`, or `type gcc`, or `type clang`. You should get output like

```
cc is /usr/bin/cc
```

If you need to install a C compiler, there are several alternatives. On a Mac, you can install Xcode from the App Store. Or you can use Homebrew to install clang or gcc:

```
brew install clang
```

or

```
brew install gcc
```

On Linux, the command would be

```
sudo apt-get install clang
```

or

```
sudo apt-get install gcc
```

2.4 Install "make" if necessary

You will also need the program "make," so type "type make." If you need to install, then

```
brew install make
```

2.5 Set up a "bin" directory to hold executable files.

On Unix-like operating systems (MacOS and Linux) it is conventional for each user to maintain a directory named "bin," just below the home directory, which contains executable files. This directory must also be added to the system's PATH variable, which is used for finding executable files.

First create a "bin" directory, if you don't already have one. To do so, first type the command

```
cd
```

3

at the shell prompt. This moves you into your home directory. Then type

```
mkdir bin
```

This creates a new directory called "bin."

You now need to add this to your shell's PATH variable. I'll assume you're using the bash shell. In your home directory, look for a file called either ".bash_profile" or ".profile". Because the name begins with ".", it will not show up if you type `ls`. However, it will appear if you type `ls .bash_profile` or `ls .profile`. If only one of these files exists, open it with a text editor (not a word processor). If they both exist, edit ".bash_profile". If neither exists, use the editor to create a new file called ".bash_profile".

Within this file, you may find existing definitions of the PATH variable. Add the following after the last line that changes this variable:

```
export PATH=$HOME/bin:$PATH
```

Save this change and exit the editor.

The contents of this file are executed by the shell each time you log into your computer. Since you have just created the file, your PATH has not yet been reset. Log out and then log back in again, and your PATH should be set. To check it, type

```
echo $PATH
```

This will print your PATH.

2.6   Compile and install legofit

The details are in the Legofit documentation.

3   Setting up your environment at the CHPC

These instructions are needed only if you're using the facilities of the Center for High Performance Computing (CHPC) at the University of Utah. We are in the process of getting that set up for the workshop. The current instructions may need to change once we get that set up.

Utah's CHPC has several "clusters," each with many "nodes." Each node is a powerful computer in its own right. The SLURM version of the pipeline submits legofit jobs to multiple nodes so that they can run in parallel. Within each node, legofit parallelizes across all available cores. This greatly reduces the time required for an analysis.

To log into a CHPC cluster, your own computer must support ssh. To find out whether it does, type the command

```
type ssh
```

at the bash prompt. If ssh is on your system, the response will be something like

```
ssh is /usr/bin/ssh
```

If ssh is not on your system, the response will be

```
-bash: type: ssh: not found
```

In the latter case, install openssh by typing

4

```
brew install openssh
```

Now that you have access to ssh, you can log into one of the CHPC clusters by typing the following at the bash prompt. Either

```
ssh -l <your unid> notchpeak.chpc.utah.edu
```

or

```
ssh -l <your unid> kingspeak.chpc.utah.edu
```

or

```
ssh -l <your unid> lonepeak.chpc.utah.edu
```

Here, `<your unid>` is your own University of Utah id. Each of these commands will log you onto one of the clusters at Utah's CHPC. Then if you haven't already done so, you'll then need to set up your CHPC environment as described above in section 3.

In the sections that follow, I will introduce you to several directories and files on the CHPC servers. Each one is described by its "path name" relative to your own "home directory," which is represented by the symbol `~`. (In scripts, this directory is often called `$HOME`.) Before these path names will work, you'll need to define two "soft links" within your own home directory, which point to group storage devices owned by the Rogers lab. After you define these soft links, `~/grp1` will refer to one of these group storage devices and `~/grp2` to the other. You can establish these links by typing

```
ln -s /uufs/chpc.utah.edu/common/home/rogersa-group1 ~/grp1
ln -s /uufs/chpc.utah.edu/common/home/rogersa-group2 ~/grp2
```

If everyone defines these soft links in the same way, then we can all use path names of the form `~/grp1/rogers/data/whatever`.

The Legofit executables are already available directory `~/grp1/bin`. To make these easily available, use a text editor to edit the file `.bash_profile` in your home directory. Within that file, you will find a line that defines a variable called `PATH`, which is used by bash to find executable programs. We want to add `~/grp1/bin` to the end of `PATH`, so that bash can find the Legofit executables. Edit that line that defines `PATH` so that it looks like this

```
export PATH=$PATH:$HOME/bin:$HOME/.local/bin:$HOME/grp1/bin
```

Also add the following line at the end of the file

```
module --latest load git
```

This last line will give you access to the program "git," which you'll need to access the agar22 repository. After making these changes, save the file and exit the editor. The file `.bash_profile` is normally read only once, when you first log in. You've only just edited this file, bash doesn't yet know about the changes you made there. To fix this, you can either type `logout` and then `ssh` back in again, or type `. .bash_profile`, which tells bash to execute the commands in this file.

Finally, clone the agar22 repository into your own home directory on the CHPC.

```
git clone https://github.com/alanrogers/agar22.git agar22
```

This will create a subdirectory called "agar22," which contains files related to the workshop.

4   Making an input file in .raf format

Legofit uses data in ".raf" format. The following data sets are already available on the server at Utah's CHPC.

Altai Neanderthal `~/grp1/rogers/data/altai/orig2/altai.raf.gz`

Vindija Neanderthal `~/grp1/rogers/data/vindija/orig/vindija.raf.gz`

Denisovan `~/grp1/rogers/data/denisova/orig2/denisova.raf.gz`

Western Europeans (SGDP) `~/grp1/rogers/data/simons/raf/weur.raf.gz`

You can use any or all of these in your projects. But you'll also need to make at least one more. To do so, copy the following script into a directory of your own:

`~/grp1/rogers/data/simons/raf/weur.slr`

Change its name (because "weur" stands for "western European") and edit it to reflect the population or populations that you want to study. It's set up to run on the "notchpeak" cluster, where I have only one node.

   If that node is free, you can run it there by using "cd" to move to the directory that contains the script and then typing

```
sbatch <your script name>
```

where `<your script name>` is the name of your version of the script. This script will take many hours to complete. You can check on its status at any time by typing

```
squeue -u <your user id>
```

where `<your user id>` is the id you use to log onto the CHPC cluster. As a shorthand, I often type this as

```
squeue -u `whoami`
```

which uses Linux's "whoami" command to fill in your user id.

   If the notchpeak node is busy, try kingspeak instead. I have six nodes there, so your odds are better. Before doing so, edit your slurm script. The lines that read

```
#SBATCH --account=rogersa-np
#SBATCH --partition=rogersa-np
```

work on notchpeak but not on kingspeak. For kingspeak, they should read

```
#SBATCH --account=rogersa-kp
#SBATCH --partition=rogersa-kp
```

5 Tabulating site patterns

Legofit works with the frequencies of "nucleotide site patterns," which are explained in this section of the Legofit website. During the workshop, you will not need to tabulate site patterns, because I'll do that for you before the workshop begins. This section is here to help you do the job on your own, after the workshop.

The Legofit package includes several programs for tabulating site patterns. I'll focus here on "sitepat," which works with data files rather than with the output of a simulation program. Sitepat reads a series of files in .raf format and writes a file describing the frequency of each site pattern. Optionally, it also writes a series of output files, each describing site pattern frequencies in a different bootstrap replicate. These replicates are used for measuring statistical uncertainty. Details are in the Legofit docs. Sitepat reads enormous data files and can take hours to run, even on a compute cluster. That is why we aren't doing this step during the workshop.

Have a look at file data.opf. After the header, the first few lines look like

```
#       SitePat             E[count]             loBnd             hiBnd
            x       312922.8110121   311974.2398439   313732.1909228
            y       306301.1413693   305178.5610866   307121.6833708
            v       102039.5699405   101293.0904018   102887.7097098
            a        78653.9538690    77779.2808408    79467.4795387
            d       339489.6889881   337728.6566221   341531.9807664
          x:y       189532.2514885   187871.4715776   191321.2242190
```

The left column lists the site patterns. "E[count]" is (more or less) the number of sites exhibiting each site pattern. (For a full explanation, see the Legofit docs.) The values under "loBnd" and "hiBnd" enclose a 95% confidence interval. The frequency of site pattern "x" is its value of "E[count]" divided by the sum of all the values under "E[count]."

It's best to run sitepat within a script, both because it takes awhile to run, and also because the script will document what you did. Here is the slurm script that generated the output above.

6 Fitting models of history

6.1 Organizing the data directory

For each set of site pattern data, I create a directory tree with the same format, which is illustrated in the github repository at legofit/europe. Within this directory, you'll find (1) a README.md file, which describes all the others, (2) a file called "data.opf," which contains observed site pattern frequencies, (3) sitepat.slr, a slurm script that runs "sitepat" and generates "data.opf," and "boot," a directory containing bootstrap replicates. There will eventually be other files with names like "all.bootci" and "all.bma," which contain the results of analyses. Within this directory, you will also find subdirectories with names like "a," "ab," and so on. Each of these refers to a different model of history, which we wish to fit to the data in "data.opf." Name these as you wish, but it's a good idea to keep the names short. In my naming scheme, directory "a" contains a model in which there is only one episode of admixture, labeled α; "ab" is for a model that has episodes α and β; and so on.

6.2 Preliminary exercises with legosim and legofit

Within the subdirectory for each model, assumptions about population history are described in file "a.lgo." These assumptions are described using syntax that you can read about here. Use `cd` to

move into directory "a," and have a look at a.lgo. Then type this at the bash command line:

```
legosim a.lgo
```

This runs the program "legosim," which reads file a.lgo, estimates the probability of each site pattern, and prints out the result. By default, legosim doesn't calculate probabilities for singleton site patterns (those in which the derived allele is present only once), and it uses a stochastic algorithm to estimate probabilities. To include singleton site patterns and use the (faster and more accurate) deterministic algorithm, the command would be

```
legosim -1 -d 0 a.lgo
```

Because it is fast, the legosim program is a useful way to check for errors in your .lgo file. For further details about the program and its command-line arguments, see its documentation.

Having ascertained that the sotware can read "a.lgo" without errors, let's use the "legofit" program to estimate parameters. The parameters to be estimated are those defined as "free" in file "a.lgo." Type the following at the bash command line:

```
legofit -1 -d 0 --tol 1e-3 a.lgo ../data.opf
```

After a few seconds, legofit should print a page of output. But before examining that, let's unpack the command that generated it. The 1st argument on the command line is "-1." This tells legofit to use singleton site patterns—those in which the derived allele is present only once. Next, the arguments "-d 0" tells legofit to use its deterministic algorithm, which is faster and more accurate than the stochastic one it uses by default. (The stochastic algorithm is needed for complex models.) The arguments "--tol 1e-3" tell legofit to be satisfied with a fairly loose fit between model and data. In research, you would want a smaller number here. The next argument, a.lgo, is the file describing the model of history, and the final argument, ../data.opf, is the data file.

Now lets examine the output of that last command. The first section of output echoes the the legofit command and the input parameters, along with default values of parameters that we didn't set on the command line.

```
#########################################
# legofit: estimate population history #
#       version 2.3.8-12-gf1f00c9       #
#########################################
#
# Program was compiled: Jun 21 2022 11:28:40
# Program was run: Thu Jul  7 20:38:55 2022
#
# cmd: legofit -1 -d 0 --tol 1e-3 a.lgo ../data.opf
# curr dir: /Users/rogers/txt/agar22/legofit/europe/a
# Stage nOptItr nSimReps
#      0    1000         1
# algorithm          : deterministic
# ignoring probs <=  : 0
# Branch length floor: 0
# DE strategy        : 2
#     F              : 0.3
#     CR             : 0.8
```

```
#     tolerance      : 0.001
# nthreads          : 2
# lgo input file     : a.lgo
# site pat input file: ../data.opf
# free parameters    : 12
# points in DE swarm : 120
# Including singleton site patterns.
# cost function      : KL
```

In the output above, "tolerance" is the value we specified as `--tol 1e-3` on the command line. The next section prints the initial values of all parameters, as read from the .lgo file.

```
Initial parameter values
Fixed:
      zero = 0
       one = 1
       TmN = 1
     Txynd = 25920
Free:
       Tnd = 24000
       Tav = 14000
       Txy = 10000
        Td = 2500
        Ta = 4000
        Tv = 1000
    twoNav = 2000
     twoNn = 2000
    twoNnd = 2000
    twoNxy = 20000
  twoNxynd = 20000
        mN = 0.01
```

Then comes a line that tells us the result of the optimization procedure.

`DiffEv reached_goal. cost=7.22802e-04 spread=8.77863e-04`

To understand this line, you need to know how legofit fits parameters to data. It uses something called "KL divergence" [5] to measure the difference between observed site pattern frequencies and those predicted by the model of history. Legofit searches for parameter values that minimize KL divergence, which is called "cost" in the output above. The value printed there (7.22802e-04) is the lowest KL divergence legofit was able to find. This search is conducted by an algorithm called "differential evolution" (DE) [7]. DE maintains a swarm of points, each representing a guess about parameter values. In each iteration of the algorithm, the best points (those with lowest KL) mutate, recombine, and reproduce to produce a new generation of points. The algorithm stops when the difference between the best point and the worst (called "spread" in the output above) falls to a specified tolerance. As you can see in this output, spread is smaller than 0.001, the tolerance we specified using "`--tol 1e-3`" on the command line. Thus, the algorithm has reached its goal, and this is why it printed `DiffEv reached_goal` in the output above. This goal is arbitrary, and the output can be useful even when the goal is not reached.

The next section of output lists the fitted values of all parameters.

```
Fitted parameter values
Free:
        Tnd = 24871.7
        Tav = 16727.1
        Txy = 13090.4
         Td = 1494.04
         Ta = 5346.73
         Tv = 2660.19
     twoNav = 14309.4
      twoNn = 8788.21
     twoNnd = 2998.19
     twoNxy = 23935.7
   twoNxynd = 46699.2
         mN = 0.0198353
```

Finally, legofit prints the expected branch length associated with each site pattern in the fitted model.

```
#       SitePat   BranchLen
              x 38204.0073508
              y 37200.1770458
              v 12628.3146863
              a 10003.0899715
              d 41976.1547361
            x:y 22591.0958761
            x:v 274.3861485
            x:a 278.8300856
            x:d 3121.4558596
            y:v 383.1800171
            y:a 321.8630482
            y:d 2936.3292837
            v:a 28571.5881777
            v:d 1060.8437350
            a:d 1067.6185093
          x:y:v 576.2188433
          x:y:a 571.7749063
          x:y:d 7001.3957667
          x:v:a 2884.5519560
          x:v:d 225.8439349
          x:a:d 230.2878720
          y:v:a 3460.0030731
          y:v:d 235.9020933
          y:a:d 229.1273190
          v:a:d 17599.3335385
        x:y:v:a 7238.2996703
        x:y:v:d 624.7610570
        x:y:a:d 620.3171199
        x:v:a:d 5567.0327885
        y:v:a:d 6019.8141157
```

For example, the expected branch length of site pattern x is 38204.0073508 generations. This is the average length, within the gene genealogy, of the branch that, should a mutation strike it, would generate site pattern x. Under the model of infinite sites [4], the expected frequency of a site pattern among polymorphic sites is proportional to its expected length and can be calculated as the ratio of this length to the sum of all the lengths.

6.3   A legofit pipeline

A minimal legofit analysis would consist of just one run of the program, such as the one we did above. This provides a point estimate of each parameter but does not estimate the uncertainties of those estimates. To estimate uncertainties, you also need to run legofit on each bootstrap estimate. Having done all that, you have completed what I like to call "stage 1" of a complete analysis.

The second stage of analysis is designed to guard against a problem that often arises when one tries to find the minimum of a complex function. The cost function may have many local minima (low spots), some of which are deeper than others. We are looking for the "global minimum," which is the deepest of these local minima. The DE algorithm is pretty good at finding global minima, but it isn't perfect. It's possible that the legofit runs from stage 1 ended up at several local minima. Stage 2 is designed to choose among these local minima.

To make this possible, each job in stage 1 uses the `--stateOut` option, which tells legofit to write a file summarizing the final state of the DE swarm. Then, each job in stage 2 uses the `--stateIn` option, which tells legofit to initialize its swarm of DE points by reading all the state files produced by stage 1. Thus, each legofit job in stage 2 begins with a DE swarm that includes points from all the legofit jobs of stage 1. If the initial swarm includes points from several local minima, then legofit can choose among these to find the global minimum.

In addition to multiple local minima, legofit also faces another problem: a small increase in the value of one parameter may have nearly the same effect as a small decrease in another. This problem can be seen in Fig. 1, where some parameter estimates are tightly correlated with others. These correlations make it hard for the optimizer to zero in on optimal parameter values.

To ameliorate this problem, I use the "pclgo" program (within the Legofit package). It uses principal components analysis to reexpress all free parameters as functions of a set of uncorrelated "principal components." Optionally, you can tell pclgo to ignore principal components that account only for a tiny fraction of the variance. This step rewrites a.lgo to produce a new file called "b.lgo." In this new file, all free variables are principal components, and all the real variables are functions of the principal components. The third and fourth stages of analysis are just like the first and second, except that I use b.lgo instead of a.lgo.

There are several ways to organize this four-stage analysis into a pipeline. On the CHPC cluster, I use SLURM scripts, which are discussed below. First, however, I will introduce a version of the pipeline that is written entirely in bash.

6.3.1   Bash version of pipeline

Within the agar22 directory tree, `cd` into directory `legofit/europe/a`. You'll find several bash scripts there, whose names all end with `.sh`. All but one of these scripts is involved in the bash version of the pipeline. To experiment with this pipeline, let's copy these scripts into a temporary directory. From within this directory, type

```
mkdir ../tmp    # make a temporary directory
cp a.lgo ../tmp # copy .lgo file into tmp
cp *.sh ../tmp  # copy scripts into tmp
```
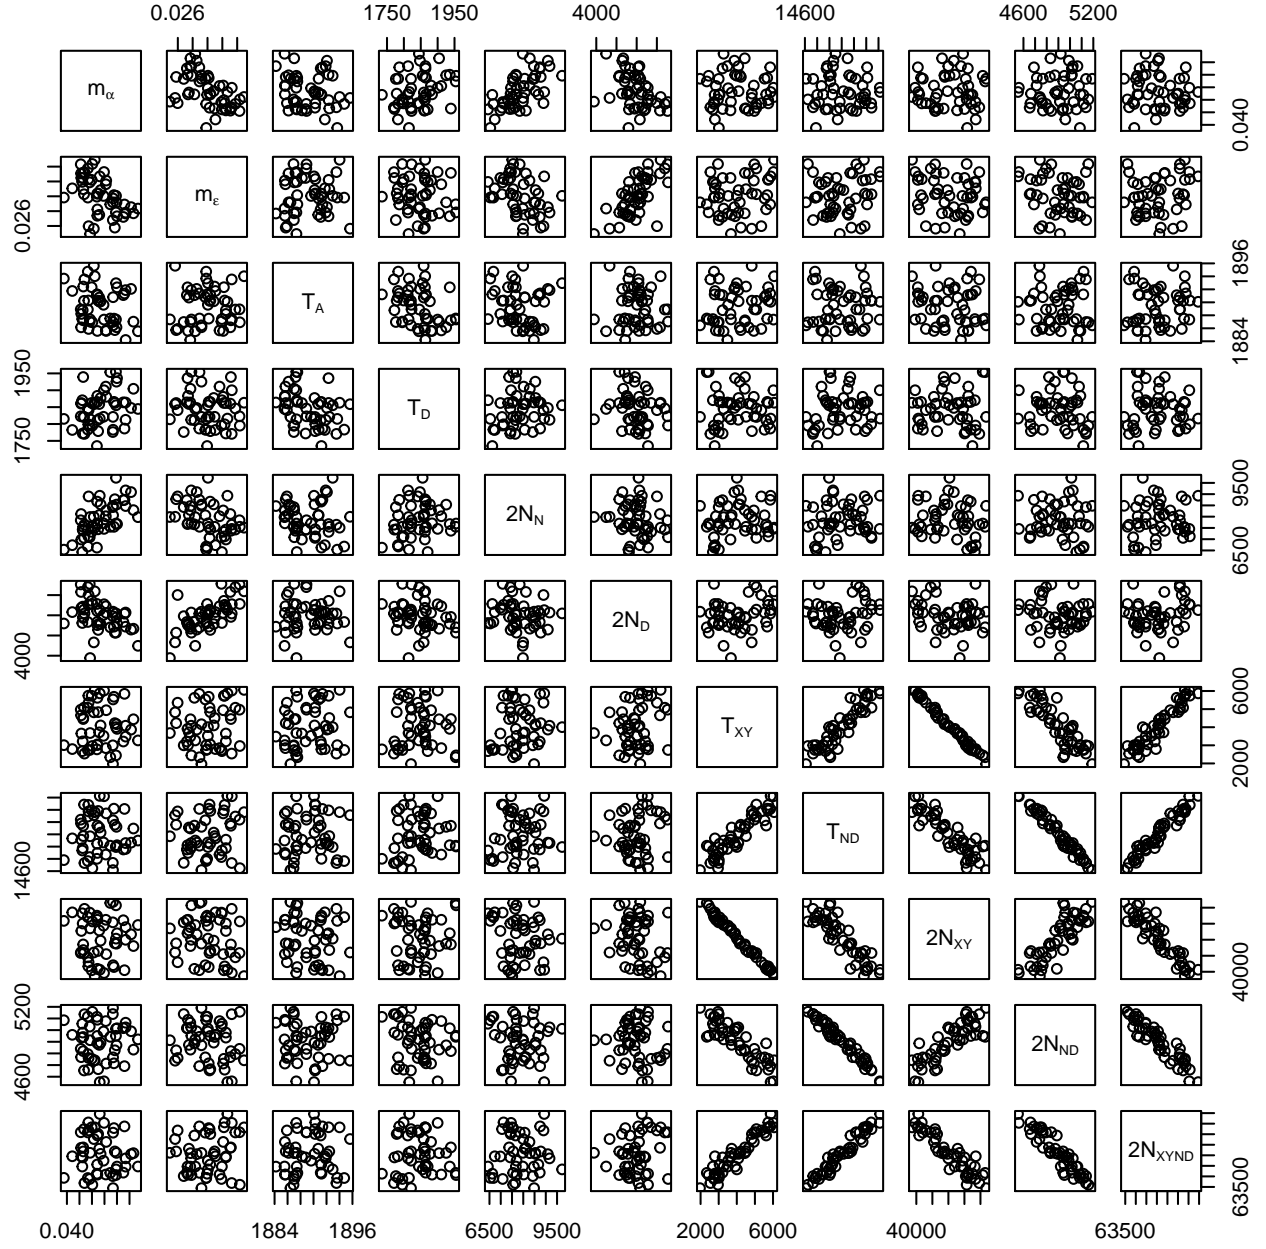
Figure 1: Scatter plot of each parameter against each other, based on 50 simulated data sets. After Rogers [9, Fig. 3].

```
cd ../tmp        # move into that directory
rm pipeline.sh   # get rid of the file that isn't part of the bash pipeline
```

Your `tmp` directory should now contain the following files:

a.lgo	defines model of history
a1.sh	stage 1 of analysis on real data
a1boot.sh	stage 1 on a bootstrap replicate
a2.sh	stage 2 on real data
a2boot.sh	stage 2 on a bootstrap replicate
pclgo.sh	run pclgo to create b.lgo
b1.sh	stage 3 of analysis on real data
b1boot.sh	stage 3 on a bootstrap replicate
b2.sh	stage 4 on real data
b2boot.sh	stage 4 on a bootstrap replicate
bash_pipeline.sh	runs all the other scripts

If you examine these scripts (for example, by typing `less a1.sh`), you'll find that the legofit commands include the option `-S 5000`, which tells the DE algorithm to do as many as 5000 iterations. This could take awhile, so let's speed things up. To change "5000" into "100" in all the files at once, type

```
sed -i "" 's/5000/100/g' *.sh
```

This will reduce the accuracy of our answers, but that's okay—this is only an exercise. To run the first stage of analysis on the real data, enter the following command:

```
bash a1.sh
```

That command takes 2.7 seconds on my 2018 Macbook Air. After it runs, there should be three new files in your directory:

a1.legofit	main output from legofit
a1.err	error output from legofit
a1.state	state of DE swarm at end of legofit run

To run the first stage on all bootstrap replicates, type:

```
seq 0 49 | xargs -n 1 bash a1boot.sh
```

This command will run 50 jobs, each of which might have been launched by typing on the command line something like

```
bash a1boot.sh <i>
```

where `<i>` is 0 for the first job, 1 for the second, and so on up to 49. Each job will analyze a single bootstrap replicate and will write three files that are analogous to the three discussed above. It takes several minutes to do all 50 bootstrap replicates.

We could, if we wished, proceed in this fashion and execute the remaining stages of analysis by typing bash commands at the terminal. But there is an easier way. Examine the file `bash_pipeline.sh` by typing

```
less bash_pipeline.sh
```

After a few lines at the top, which establish rules for dealing with errors, you will find

```
# Stage 1 of analysis.
bash a1.sh
seq 0 49 | xargs -n 1 bash a1boot.sh
```

These lines are identical to the commands we used above to execute the first stage of analysis. The script as a whole does the entire legofit analysis. To execute it all, you would type

```
bash bash_pipeline.sh
```

You would end up with 204 .legofit files (51 for each stage), the same number of .err files, and half that many .state files. Before talking about what to do with these files, let us consider other ways to implement a pipeline.

This implementation, using bash, has the advantages of simplicity and low overhead. It's probably the easiest version to get working on your laptop. However, there are better ways to implement a pipeline. One problem the current approach is that when something goes wrong half way through the analysis, it's hard to start the process up again in the middle. Another is that we are doing only one legofit run at a time. It would be better to do many runs in parallel.

In the bash pipeline, we used the xargs command. This has a "-P" argument, which makes it possible to run multipe jobs in parallel. This would not be useful in the current pipeline, however, because legofit is multithreaded and each run uses all the cores available on your computer. To run multiple legofit jobs in parallel, you need a compute cluster such as those available at the University of Utah's Center for High Performance Computing (CHPC).

6.3.2 SLURM version of pipeline

Use ssh as described above to log into one of the CHPC clusters, and then cd into directory agar22/legofit/europe/a. You'll find several SLURM scripts there, whose names all end with .slr. There is also a shell script called "pipeline.sh," which runs the entire pipeline. To experiment with this pipeline, let's copy these scripts into a temporary directory. From within this directory, type

```
mkdir ../tmp          # make a temporary directory
cp a.lgo ../tmp       # copy .lgo file into tmp
cp *.slr ../tmp       # copy slurm scripts into tmp
cp pipeline.sh ../tmp # copy pipeline script
cd ../tmp             # move into that directory
```

Your tmp directory should now contain the following files:

a.lgo	defines model of history
a1.slr	stage 1 of analysis on real data
a1boot.slr	stage 1 on a bootstrap replicate
a2.slr	stage 2 on real data
a2boot.slr	stage 2 on a bootstrap replicate
pclgo.slr	run pclgo to create b.lgo
b1.slr	stage 3 of analysis on real data
b1boot.slr	stage 3 on a bootstrap replicate
b2.slr	stage 4 on real data
b2boot.slr	stage 4 on a bootstrap replicate
pipeline.sh	runs all the SLURM scripts

To speed things up, we can reduce the number of iterations (as explained above) by typing

```
sed -i "" 's/5000/100/g' *.slr
```

Now each legofit job will do only 100 DE iterations. The results won't be as precise, but the pipeline will be fast. To run the first stage of analysis on the real data, enter the following command

```
sbatch -A rogersa -p <your cluster> a1.slr
```

where `<your cluster>` is "notchpeak," "kingspeak," or "lonepeak," depending on which cluster you logged into. This queues your job for execution but does not necessarily run it immediately. If the system is busy, it may not run for hours. To check on the status of this job, type

```
squeue -u <your unid>
```

The output will summarize the status of all jobs you've submitted to SLURM.

To run the first stage of analysis on all bootstrap replicates, the command is

```
sbatch -A rogersa -p <your cluster> --array=0-49 a1boot.slr
```

To launch the entire pipeline, first use a text editor (e.g., vim or emacs) to edit the file pipeline.sh. Near the top of that file, you'll find the lines

```
export SBATCH_ACCOUNT=rogersa
export SBATCH_PARTITION=notchpeak
```

Edit these lines to reflect the account and partition you are using, then save the file and exit the editor. Finally, type

```
. pipeline.sh
```

This will queue the entire pipeline for execution.

When the pipeline finishes, you'll want to copy the results back to your own machine for analysis. One way to do this uses the "rsync" command, which you may need to install on your machine by typing

```
brew install rsync
```

Then, to get the results of the SLURM pipeline, type this from your own computer:

```
rsync -Cavzbu <your_unid>@notchpeak.chpc.utah.edu:agar22/legofit/europe/tmp .
```

This assumes that the results you want are on the CHPC server in a directory called `~/agar22/legofit/europe/tmp`. It will create on your own machine a directory called "tmp" in the directory from which you issued the rsync command. That directory should now contain all the results from the SLURM pipeline.

Incidentally, rsync is a fairly awkward tool for copying things back and forth between the CHPC server and your own machine. I usually do this by setting up a repository on github, which I can access both from the CHPC and from my home machine. To transfer from CHPC to home, I do `git push` from the CHPC and then `git pull` from home. I'm not teaching you how to do this here, because this tutorial is complicated enough already. I'm teaching the rsync approach instead, because it's simpler.

6.3.3 Snakemake version of the pipeline

This section hasn't been written yet.

7 Model selection and model averaging

As discussed above, legofit uses KL divergence to measure the difference between observed and expected site pattern frequences, and it labels this value "cost" in its output. One might therefore assume that the best model would be the one with the lowest value of cost. This is not necessarily so, because of the problem of overfitting. Very complex models fit not only the signal in the data, but also the noise, and the noise in one data set does not predict that in another. Consequently, the best-fitting model may not be the most predictive.

To address this problem, Legofit uses the "bootstrap estimate of predictive error" (bepe) [8, 2, 3]. This method uses variation among data sets (the real data plus 50 replicates generated by a moving-blocks bootstrap [6]) to approximate variation in repeated sampling. It fits the model to one data set and then tests this fit against all the others. Let us apply Legofit's "bepe" program to the output from the last stage of analysis. Use `cd` to move into the directory with all your .legofit output files, and type

```
bepe ../data.opf ../boot/boot*.opf -L b2.legofit b2boot*.legofit > b2.bepe
```

On the bepe command line, we first list all the input .opf files, beginning with the real data. Then, after `-L` we list the corresponding .legofit output files in the same order. The first few lines of output look like this

```
####################################################
# bepe: bootstrap estimate of predictive error #
#           version 2.3.8-12-gf1f00c9           #
####################################################
#
# Program was compiled: Mar 22 2022 17:40:38
# Program was run: Sat Jul  9 13:41:24 2022
# curr dir: /Users/rogers/txt/agar22/legofit/europe/tmp
#
#          bepe         DataFile      LegofitFile
1.130271806e-06          data.opf        b2.legofit
1.089578158e-06          boot0.opf  b2boot0.legofit
```

The first number in the left column is the bepe value of the real data. This is the number to use in choosing among models. The numbers below this one refer to bootstrap replicates and are use by booma, to which we now turn.

If one model has a bepe value much lower than the others, then it's a pretty obvious winner. But what if the best bepe is only slightly lower than the second best? When several models provide reasonable descriptions of the data, it is better to average across models than to choose just one. This allows uncertainty about the model itself to be incorporated into confidence intervals. For this purpose, Legofit uses bootstrap model averaging, "booma" [1, 8]. The booma weight of the ith model is the fraction of data sets (including the real data and 50 bootstrap replicates) in which that model "wins," i.e. has the lowest value of bepe.

If you have fit several models to the same data, you can use Legofit's "booma" program to calculate model-averaged estimates. Before running booma, make sure that the directory for each

model has a file named "b2.bepe." Each directory should also contain a file named "b2.flat," which contains a row for each data set and a column for each estimated parameter. To create these files, `cd` into the directory for each model, and type

```
flatfile.py b2.legofit b2boot*.legofit > b2.flat
```

Within the agar22 repository, all the model directories with `legofit/europe` already have both b2.bepe and b2.flat. From the parent directory (the one that contains data.opf), type

```
booma */b2.bepe -F */b2.flat > all.bma
```

This will create a file called "all.bma" that gives the booma weight of each model and also the model-averaged parameter estimates. The first few lines look like this

```
####################################
# booma: bootstrap model average #
#    version 2.3.9-7-g8bbcef2     #
####################################
#
# Program was compiled: Jul  9 2022 14:41:56
# Program was run: Sat Jul  9 14:50:30 2022
# cmd: booma a/b2.bepe ab/b2.bepe abc/b2.bepe abcd/b2.bepe abd/b2.bepe ac/b2.be>
# curr dir: /Users/rogers/txt/agar22/legofit/europe
#         Weight         MSC_file        Flat_file
#              0          a/b2.bepe        a/b2.flat
#              0         ab/b2.bepe       ab/b2.flat
#              0        abc/b2.bepe      abc/b2.flat
#    0.9803921569      abcd/b2.bepe     abcd/b2.flat
#   0.01960784314       abd/b2.bepe      abd/b2.flat
#              0         ac/b2.bepe       ac/b2.flat
#              0        acd/b2.bepe      acd/b2.flat
#              0         ad/b2.bepe       ad/b2.flat
```

Only two of the models (abc and abcd) have non-zero weights. These weights are used in calculated the model-averaged parameter estimates, which are listed below the weights.

To calculate the bootstrap confidence interval of each parameter, type

```
bootci.py all.bma > all.bootci
```

These confidence intervals are slightly wider than they would be, had we based them on results from a single model, rather than on an average across models. This is because they include uncertainty about the model itself in addition to uncertainty about parameter values under a given model.

8 Graphing results

References

[1]   Steven T Buckland, Kenneth P Burnham, and Nicole H Augustin. "Model Selection: an Integral Part of Inference". In: *Biometrics* 53.2 (1997), pp. 603–618. DOI: 10.2307/2533961.

[2]   Bradley Efron. "Estimating the Error Rate of a Prediction Rule: Improvement on Cross-Validation". In: *Journal of the American Statistical Association* 78.382 (1983), pp. 316–331. DOI: 10.1080/01621459.1983.10477973.

[3]   Bradley Efron and Robert J. Tibshirani. *An Introduction to the Bootstrap*. New York: Chapman and Hall, 1993. DOI: 10.1007/978-1-4899-4541-9.

[4]   Motoo Kimura. "The Number of Heterozygous Nucleotide Sites Maintained in a Finite Population Due to Steady Flux of Mutation". In: *Genetics* 61 (1969), pp. 893–903. DOI: 10.1093/genetics/61.4.893.

[5]   Solomon Kullback and Richard A Leibler. "On Information and Sufficiency". In: *The Annals of Mathematical Statistics* 22.1 (Mar. 1951), pp. 79–86. DOI: 10.1214/aoms/1177729694.

[6]   Regina Y. Liu and Kesar Singh. "Moving Blocks Jacknife and Bootstrap Capture Weak Dependence". In: *Exploring the "Limits" of the Bootstrap*. Ed. by Raoul LePage and Lynne Billard. New York: Wiley, 1992, pp. 225–248.

[7]   Kenneth Price, Rainer M Storn, and Jouni A Lampinen. *Differential Evolution: A Practical Approach to Global Optimization*. Berlin: Springer Science and Business Media, 2006. ISBN: 978-3-540-20950-8.

[8]   Alan R. Rogers. "Legofit: Estimating Population History from Genetic Data". In: *BMC Bioinformatics* 20 (2019), p. 526. DOI: 10.1186/s12859-019-3154-1.

[9]   Alan R. Rogers. "An Efficient Algorithm for Estimating Population History from Genetic Data". In: *Peer Community Journal* 2 (2022), e32. DOI: 10.24072/pcjournal.132.