A Legofit Tutorial

Alan R. Rogers

July 5, 2022

1 Introduction

This is a tutorial on Legofit, a package that uses genetic data to estimate the history of population size, subdivision, and admixture [1, 2]. The full documentation is available online. It will step you through the process of installing the software, using it to tabulate the frequencies of nucleotide site patterns, and then analyzing these data to fit a model of population history. We will assume that data files are already available in .vcf format.

This tutorial will be used in the 2022 edition of the AGAR Workshop. During the workshop, we will do only last step in this process, which estimates population history from site pattern frequencies. We will not do the preliminary steps, which tabulate these frequencies. Nonetheless, this tutorial describes the whole process.

We plan to introduce students to the Center for High-Performance Computing (CHPC) at the University of Utah, and some of the sections below assume that you have an account on that system. These sections involve using the SLURM Workload Manager to submit jobs to the compute cluster. However, our focus during the workshop will be on the last step, which can be done in two ways: either on the CHPC (using SLURM), or on your own computer (using the bash shell). This tutorial will cover both environments.

2 Installing and Using Legofit

2.1 Install git if necessary

You may have git already. To find out whether you do, type:

```
type git
```

On Mac or Linux, this will print the location of the executable git file. If there is no such file, you'll get an error message saying `type: git: not found`. In that case, install git:

```
brew install git
```

2.2 Clone legofit

Use git to clone Legofit onto your machine. This step will create a subdirectory called "legofit," which will contain the source code. Before executing the command below, use the `cd` command to move into the directory where you keep source code that other people have written. I keep such code in a directory called `distrib`, which I originally created by typing `mkdir distrib`. Having created such a directory, move into it by typing

```
cd distrib
```

Then clone legofit by typing:

```
git clone git@github.com:alanrogers/legofit.git legofit
```

This will create a directory called legofit. Use `cd` to move into it.

2.3   Install a C compiler if necessary

To see whether you have a C compiler already, type `type cc`, or `type gcc`, or `type clang`. You should get output like

```
cc is /usr/bin/cc
```

If you need to install a C compiler, there are several alternatives. On a Mac, you can install Xcode from the App Store. Or you can use Homebrew to install clang or gcc:

```
brew install clang
```

or

```
brew install gcc
```

On Linux, the command would be

```
sudo apt-get install clang
```

or

```
sudo apt-get install gcc
```

2.4   Install "make" if necessary

You will also need the program "make," so type "type make." If you need to install, then

```
brew install make
```

2.5   Set up a "bin" directory to hold executable files.

On Unix-like operating systems (MacOS and Linux) it is conventional for each user to maintain a directory named "bin", just below the home directory, which contains executable files. This directory must also be added to the system's PATH variable, which is used for finding executable files.

First create a "bin" directory, if you don't already have one. To do so, first type the command

```
cd
```

at the shell prompt. This moves you into your home directory. Then type

```
mkdir bin
```

This creates a new directory called "bin."

You now need to add this to your shell's PATH variable. I'll assume you're using the bash shell. In your home directory, look for a file called either ".bash_profile" or ".profile". Because the name begins with ".", it will not show up if you type `ls`. However, it will appear if you type `ls .bash_profile` or `ls .profile`. If only one of these files exists, open it with a text editor (not a word processor). If they both exist, edit ".bash_profile". If neither exists, use the editor to create a new file called ".bash_profile".

Within this file, you may find existing definitions of the PATH variable. Add the following after the last line that changes this variable:

```
export PATH=$HOME/bin:$PATH
```

Save this change and exit the editor.

The contents of this file are executed by the shell each time you log into your computer. Since you have just created the file, your PATH has not yet been reset. Log out and then log back in again, and your PATH should be set. To check it, type

```
echo $PATH
```

This will print your PATH.

2.6 Compile and install legofit

The details are in the Legofit documentation.

2.7 Soft links to group storage

In the sections that follow, I will introduce you to several directories and files on the CHPC servers. Each one is described by its "path name" relative to your own "home directory," which is represented by the symbol `~`. Before these path names will work, you'll need to define two "soft links" within your own home directory, which point to group storage devices owned by the Rogers lab. After you define these soft links, `~/grp1` will refer to one of these group storage devices and `~/grp2` to the other. You can establish these links by typing

```
ln -s /uufs/chpc.utah.edu/common/home/rogersa-group1 ~/grp1
ln -s /uufs/chpc.utah.edu/common/home/rogersa-group2 ~/grp2
```

If everyone in the lab defines these soft links in the same way, then we can all use path names of the form `~/grp1/rogers/data/whatever`.

2.8 Making an input file in .raf format

Legofit uses data in ".raf" format. The following data sets are already available on the server at Utah's CHPC.

Altai Neanderthal `~/grp1/rogers/data/altai/orig2/altai.raf.gz`

Vindija Neanderthal `~/grp1/rogers/data/vindija/orig/vindija.raf.gz`

Denisovan `~/grp1/rogers/data/denisova/orig2/denisova.raf.gz`

Western Europeans (SGDP) `~/grp1/rogers/data/simons/raf/weur.raf.gz`

You can use any or all of these in your projects. But you'll also need to make at least one more. To do so, copy the following script into a directory of your own:

> `~/grp1/rogers/data/simons/raf/weur.slr`

Change its name (because "weur" stands for "western European") and edit it to reflect the population or populations that you want to study. It's set up to run on the "notchpeak" cluster, where I have only one node.

If that node is free, you can run it there by using "cd" to move to the directory that contains the script and then typing

> `sbatch <your script name>`

where `<your script name>` is the name of your version of the script. This script will take many hours to complete. You can check on its status at any time by typing

> `squeue -u <your user id>`

where `<your user id>` is the id you use to log onto the CHPC cluster. As a shorthand, I often type this as

> `squeue -u ‘whoami‘`

which uses Linux's "whoami" command to fill in your user id.

If the notchpeak node is busy, try kingspeak instead. I have six nodes there, so your odds are better. Before doing so, edit your slurm script. The lines that read

> `#SBATCH --account=rogersa-np`
> `#SBATCH --partition=rogersa-np`

work on notchpeak but not on kingspeak. For kingspeak, they should read

> `#SBATCH --account=rogersa-kp`
> `#SBATCH --partition=rogersa-kp`

2.9 Working with Legofit

2.9.1 Organizing the top-level directory for a Legofit project

For each project, I create a directory tree with the same format, which is illustrated in the github repository at legofit/europe. The top-level directory contains the following files:

README.md Describes how all the other files in the directory were made.

data.opf Observed site-pattern frequencies, as generated by Legofit's "sitepat" program, under the control of "sitepat.slr."

boot A subdirectory containing bootstrap replicates, which have names like "boot0.opf," "boot1.opf," and so on. These files are also generated by "sitepat," under the control of "sitepat.slr."

sitepat.slr A slurm script that runs "sitepat" and creates "data.opf" and all the .opf files within subdirectory "boot."

Eventually, the top-level directory will contain other files, with names like "all.bootci" and "all.bma," which contain the results of analyses.

Within the top-level directory, there should be a directory for each model that is fit to the data in "data.opf" and "boot/boot*.opf." Name these as you wish, but it's a good idea to keep the names short. Mine tend to have names like "a," "b," "ab," etc. The "a" directory contains a model in which there is only one episode of admixture, labeled α. The "ab" directory is for a model that has episodes α and β.

2.9.2   Tabulating site patterns

The next step is to run sitepat, which reads a series of .raf files and writes a file describing the frequency of each site pattern. Optionally, it also writes a series of output files, each describing site pattern frequencies in a different bootstrap replicate. These replicates are used for measuring statistical uncertainty. Details are in the Legofit docs. Sitepat reads enormous data files and can take hours to run, even on a compute cluster. In this workshop, I'll do this step for you.

Have a look at file data.opf. After the header, the first few lines look like

```
#      SitePat          E[count]            loBnd            hiBnd
            x     312922.8110121    311974.2398439    313732.1909228
            y     306301.1413693    305178.5610866    307121.6833708
            v     102039.5699405    101293.0904018    102887.7097098
            a      78653.9538690     77779.2808408     79467.4795387
            d     339489.6889881    337728.6566221    341531.9807664
          x:y     189532.2514885    187871.4715776    191321.2242190
```

The left column lists the site patterns. "E[count]" is (more or less) the number of sites exhibiting each site pattern. (For a full explanation, see the Legofit docs.)

2.9.3   Running legofit

On the repo, look at directory legofit/europe/a. This deals with a simple model, with only one episode of admixture. The model of history is described in file a.lgo, using a syntax described in the Legofit documentation. You will want to write your own .lgo file, which describes the model of history you want to study. But first, let's do an exercise with this one.

Before attempting this exercise, you should have installed Legofit as described above and cloned the agar22 repo onto your own computer. Then use the terminal application to `cd` into directory `legofit/europe/a` and type:

```
legofit -1 -d 0 --tol 1e-3 a.lgo ../data.opf
```

After a few seconds, legofit should print a page of output. But before examining that, let's unpack the command that generated it. The 1st argument on the command line is "`-1`." This tells legofit to use singleton site patterns—those in which the derived allele is present only once. Next, the arguments "`-d 0`" tells legofit to use its deterministic algorithm, which is faster and more accurate than the stochastic one it uses by default. (The stochastic algorithm is needed for complex models.) The arguments "`--tol 1e-3`" tell legofit to be satisfied with a fairly loose fit between model and data. In research, you would want a smaller number here. The next argument, `a.lgo`, is the file describing the model of history, and the final argument, `../data.opf`, is the data file.

2.9.4 Studying a model of history

For any given data set, you will want to study several models. It is convenient to keep the files relating to a given model in a directory just under the one that holds the data file. The directory for each model contains the following files:

a.lgo describes the model of history under study. The syntax of a .lgo file is described in the legofit documentation.

a1.slr is a script, which can be interpreted by "slurm," a program that manages jobs on a compute cluster. This file runs legofit on the real data.

a1boot.slr a slurm script that runs legofit on a bootstrap replicate. This script is run once for each bootstrap replicate, using slurm's `sbatch --array` command.

a2.slr

```
sed -i "" 's/5000/100/g' *.sh
```

References

[1] Alan R. Rogers. "Legofit: Estimating Population History from Genetic Data". In: *BMC Bioinformatics* 20 (2019), p. 526. DOI: `10.1186/s12859-019-3154-1`.

[2] Alan R. Rogers. "An Efficient Algorithm for Estimating Population History from Genetic Data". In: *Peer Community Journal* 2 (2022), e32. DOI: `10.24072/pcjournal.132`.